

7 steps for highly effective deep neural networks

Natalino Busa - Head of Data Science



Natalino Busa



[View Profile](#)

Linkedin and Twitter:
@natbusa

O'Reilly Author and Speaker

Teradata EMEA Practice Lead on Open Source Technologies

Teradata Principal Data Scientist

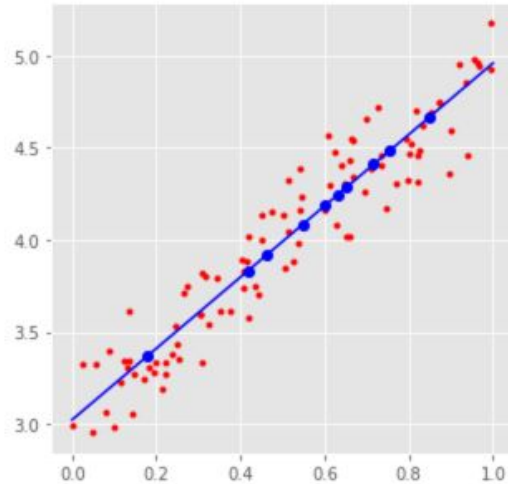
ING Group Enterprise Architect: Cybersecurity, Marketing, Fintech

Cognitive Finance Group Advisory Board Member

Philips Senior Researcher, Data Architect

Linear Regression:

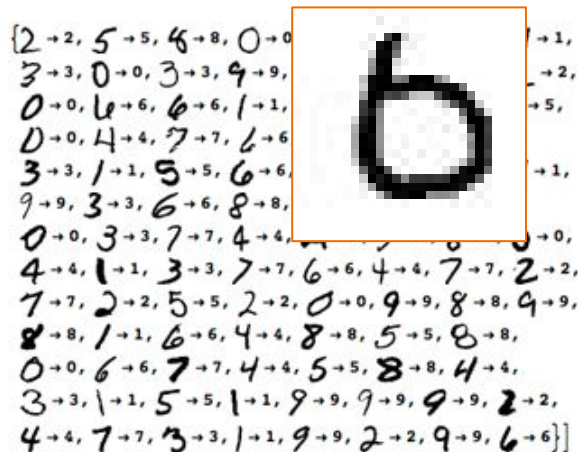
How to best fit a line to some data



X: independent variable
Y: dependent variable

Classification: Handwritten digits

28x28 pixels



input : 784 numbers

output: 10 classes

Sharing the (Not so) Secret Lore:

Keras & Tensorflow

Some smaller projects:

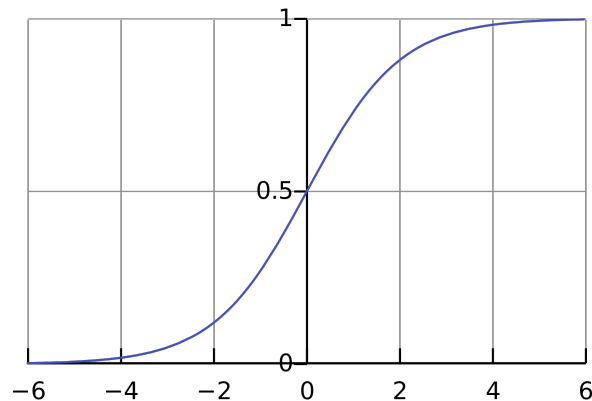
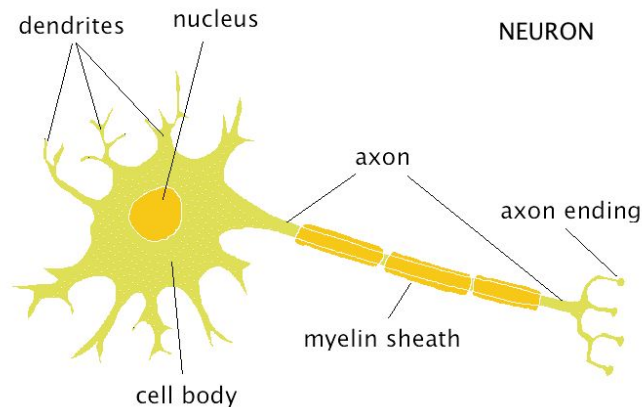
[Tflearn](#), [Tensorlayers](#)



<http://keras.io/>



1: Single Layer Perceptron



Activation function

Axon's response

$$\theta_i = \frac{1}{1 + \exp \left[- \left(\beta_0 + \sum_{j=1}^k \beta_j x_{ij} \right) \right]}$$

"dendrites"

More activation functions:

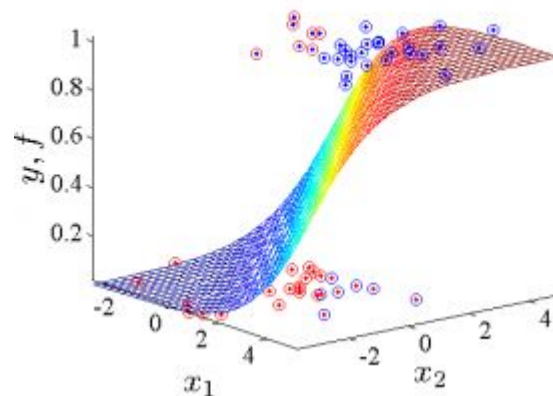
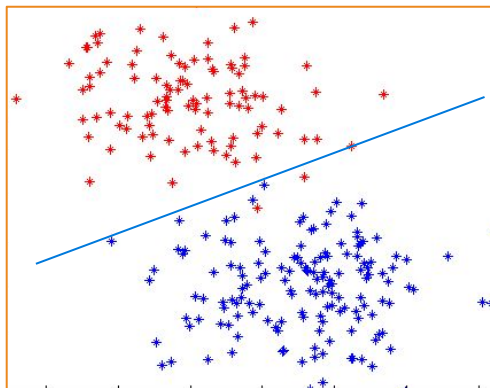
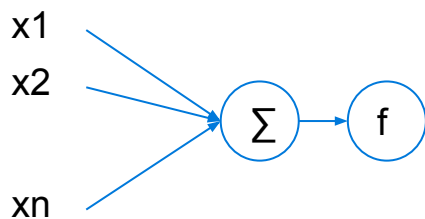
https://en.wikipedia.org/wiki/Activation_function

Name	Plot	Equation	Derivative (with respect to x)	Range	Order of continuity	Monotonic	Derivative Monotonic	Approximates identity near the origin
Identity		$f(x) = x$	$f'(x) = 1$	$(-\infty, \infty)$	C^∞	Yes	Yes	Yes
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$	$\{0, 1\}$	C^{-1}	Yes	No	No
Logistic (a.k.a. Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$	$(0, 1)$	C^∞	Yes	No	No
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$	$(-1, 1)$	C^∞	Yes	No	Yes
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$	$\left(-\frac{\pi}{2}, \frac{\pi}{2}\right)$	C^∞	Yes	No	Yes
Softsign ^{[7][8]}		$f(x) = \frac{x}{1 + x }$	$f'(x) = \frac{1}{(1 + x)^2}$	$(-1, 1)$	C^1	Yes	No	Yes
Rectified linear unit (ReLU) ^[9]		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$[0, \infty)$	C^0	Yes	Yes	No
Leaky rectified linear unit (Leaky ReLU) ^[10]		$f(x) = \begin{cases} 0.01x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0.01 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$(-\infty, \infty)$	C^0	Yes	Yes	No
Parametric rectified linear unit (PReLU) ^[11]		$f(\alpha, x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(\alpha, x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$(-\infty, \infty)$	C^0	Yes iff $\alpha \geq 0$	Yes	Yes iff $\alpha = 1$
Randomized leaky rectified linear unit (RRReLU) ^[12]		$f(\alpha, x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$ ^[1]	$f'(\alpha, x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$(-\infty, \infty)$	C^0	Yes	Yes	No
Exponential linear unit (ELU) ^[13]		$f(\alpha, x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(\alpha, x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$(-\alpha, \infty)$	C^1 when $\alpha = 1$, otherwise C^0	Yes iff $\alpha \geq 0$	Yes iff $0 \leq \alpha \leq 1$	Yes iff $\alpha = 1$

1: Single Layer Perceptron (binary classifier)

Single Layer Neural Network

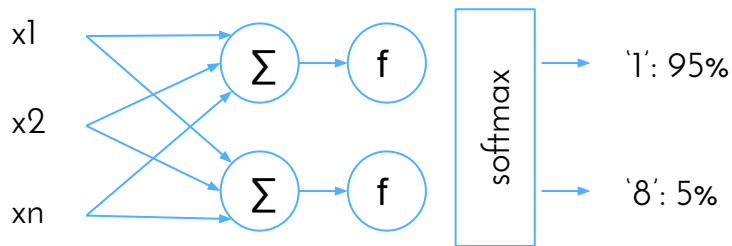
Takes: n-input features: Map them to a soft “binary” space



1: Single Layer Perceptron (multi-class classifier)

From soft binary space to predicting probabilities:

Take n inputs, Divide by the sum of the predicted values



$$P(y = j|\mathbf{x}) = \frac{e^{\mathbf{x}^T \mathbf{w}_j}}{\sum_{k=1}^K e^{\mathbf{x}^T \mathbf{w}_k}}$$

Values between 0 and 1

Sum of all outcomes = 1

It produces an **estimate of a probability!**

1: Single Layer Perceptron

Minimize costs:

The cost function depends on:

- Parameters of the model
- How the model “composes”

Goal :

Reduce the mean probability error

modify the parameters to reduce the error!

Gradient Descent Algorithm:

$$\Theta_{n+1} = \Theta_n - \alpha \frac{\partial}{\partial \Theta_n} J(\Theta_n)$$

$\Theta \rightarrow$ Parameter Vector

$J \rightarrow$ Cost Function

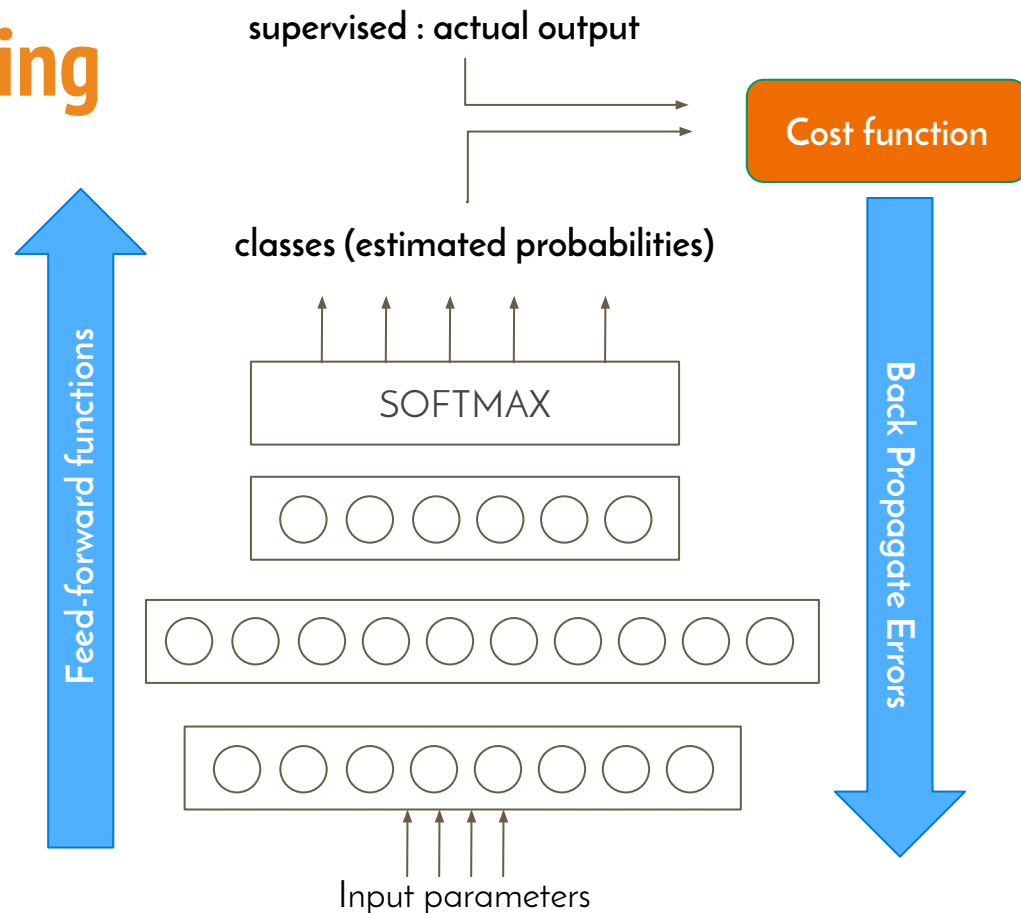
$\alpha \rightarrow$ Slope Parameter

Vintage math from last century

Supervised Learning

Stack layers of perceptrons

- Feed Forward Network
- Scoring goes from input to output
- Back propagate the error from output to input



Let's go!

```
In [7]: %matplotlib inline
```

```
import numpy as np
import matplotlib.pyplot as plt
```

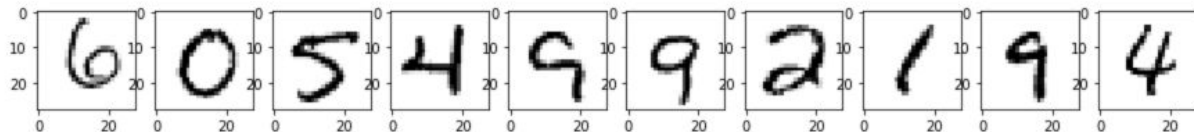
```
In [8]: from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("../mnist-data/", one_hot=True)
```

```
Extracting ../mnist-data/train-images-idx3-ubyte.gz
Extracting ../mnist-data/train-labels-idx1-ubyte.gz
Extracting ../mnist-data/t10k-images-idx3-ubyte.gz
Extracting ../mnist-data/t10k-labels-idx1-ubyte.gz
```

```
In [9]: mnist.train.images.shape
```

```
Out[9]: (55000, 784)
```

```
In [10]: plt.figure(figsize=(15,5))
for i in list(range(10)):
    plt.subplot(1, 10, i+1)
    pixels = mnist.test.images[i+100]
    pixels = pixels.reshape((28, 28))
    plt.imshow(pixels, cmap='gray_r')
plt.show()
```



1. Single Layer Perceptron

```
In [11]: from keras.models import Sequential
         from keras.layers.core import Dense, Activation
```

```
In [12]: model = Sequential()
         model.add(Dense(10, input_shape=(784,)))
         model.add(Activation('softmax'))
```

```
In [13]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [14]: model.fit(mnist.train.images, mnist.train.labels,
                   batch_size=500, nb_epoch=10, verbose=1,
                   validation_data=(mnist.test.images, mnist.test.labels))
```



1. Single Layer Perceptron

```
In [10]: score = model.evaluate(mnist.test.images, mnist.test.labels, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
Test score: 0.293662216854
Test accuracy: 0.9191
```

```
In [11]: # test item #100 is a six
pixels = mnist.test.images[100]
result = model.predict_on_batch(np.array([pixels]))
dict(zip(range(10), result[0]))
```

```
Out[11]: {0: 0.0073514683,
1: 0.0034446407,
2: 0.047751036,
3: 0.00301607,
4: 0.0032819158,
5: 0.00030964505,
6: 0.92316657,
7: 0.00083195668,
8: 0.0075355168,
9: 0.0033110771}
```

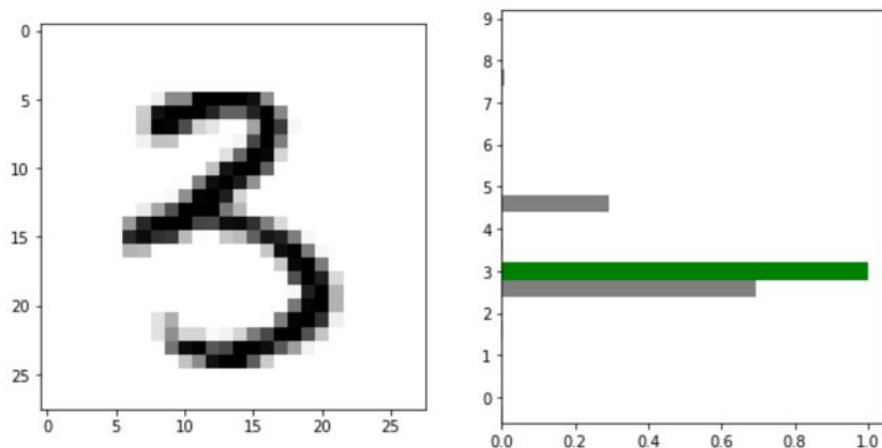


1. Single Layer Perceptron

```
In [17]: import random
i = random.randint(0,mnist.test.images.shape[0])

pixels = mnist.test.images[i]
truth = mnist.test.labels[i]
result = model.predict_on_batch(np.array([pixels]))[0]

test_render(pixels, result, truth)
```



1. Single Layer Perceptron



```
In [7]: # Set parameters
learning_rate = 0.01
training_iteration = 10
batch_size = 250
```

```
FLAGS = None
```

```
In [8]: # TF graph input
x = tf.placeholder('float', [None, 784]) # mnist data image of shape 28*28=784
y = tf.placeholder('float', [None, 10]) # 0-9 digits recognition => 10 classes
```

```
In [9]: # Set model weights
W = tf.Variable(tf.zeros([784, 10]), name='W')
b = tf.Variable(tf.zeros([10]), name='b')
```

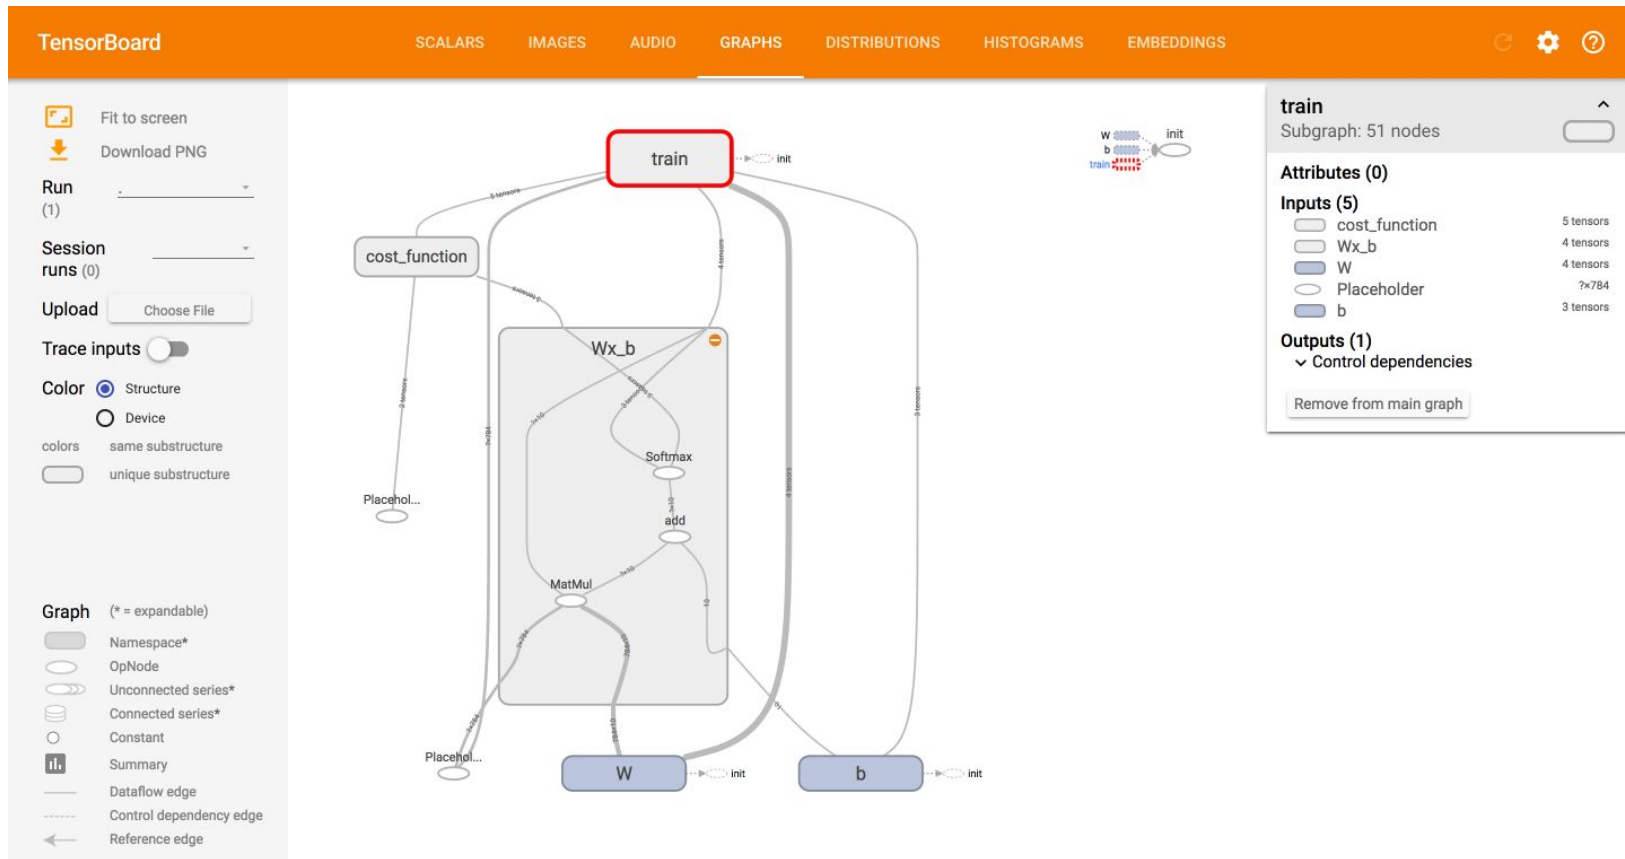
```
In [10]: with tf.name_scope("Wx_b") as scope:
# Construct a linear model
y_hat = tf.nn.softmax(tf.matmul(x, W) + b) # Softmax

# Add summary ops to collect data
tf.summary.histogram("weights", W)
tf.summary.histogram("biases", b)
```

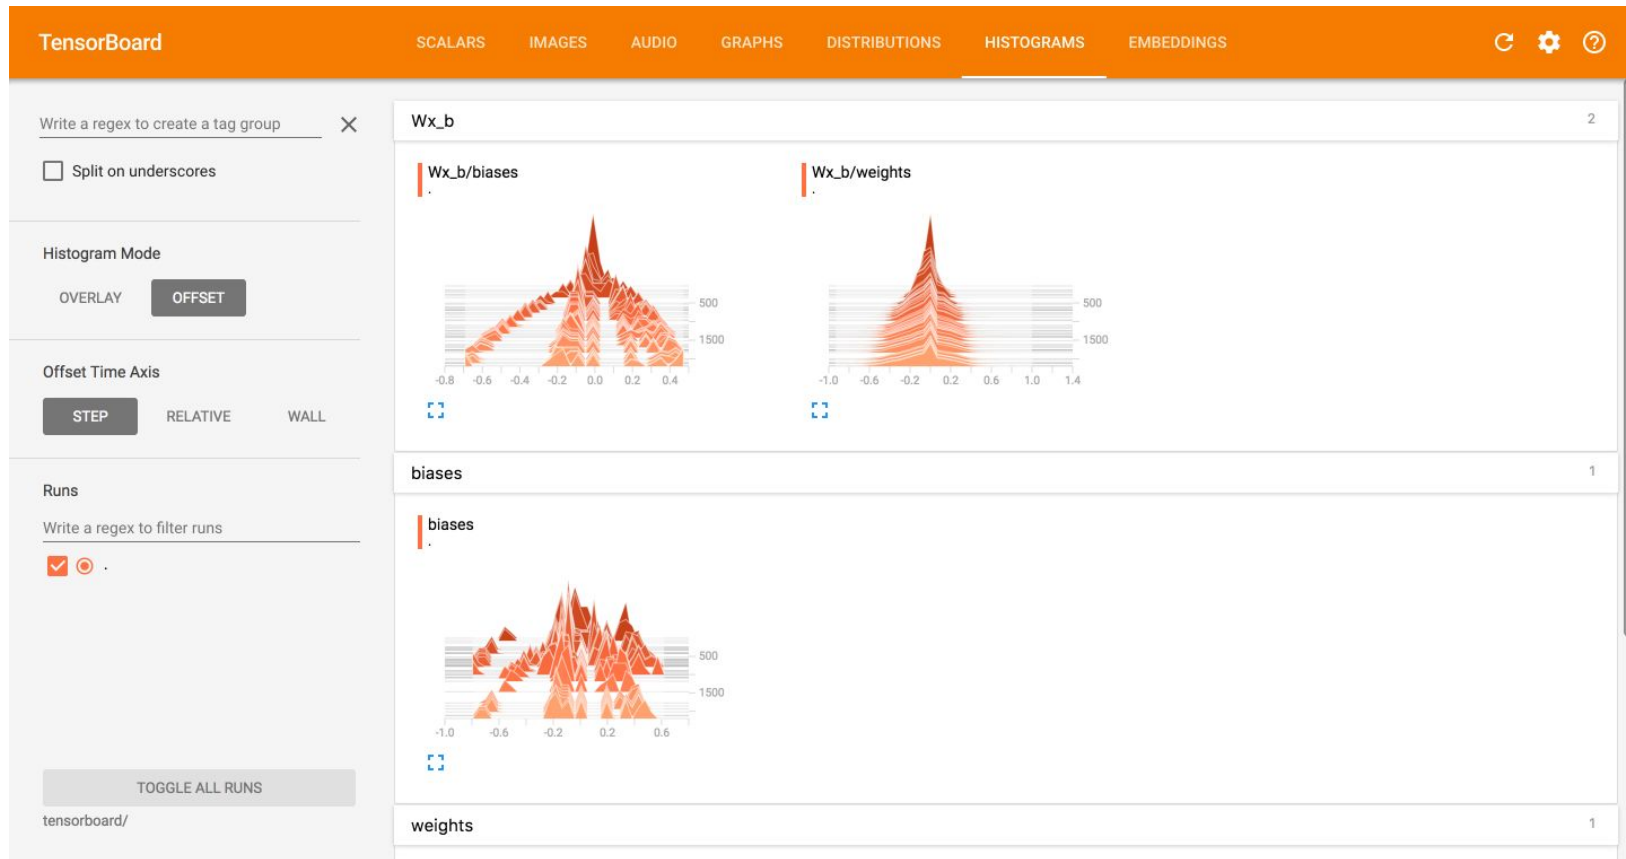
```
In [11]: # More name scopes will clean up graph representation
with tf.name_scope("cost_function") as scope:
# Minimize error using cross entropy
# Cross entropy
cost_function = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=y, logits=y_hat))
# Create a summary to monitor the cost function
tf.summary.scalar("cost_function", cost_function)
```

```
In [12]: with tf.name_scope("train") as scope:
# Gradient descent
optimizer = tf.train.AdamOptimizer().minimize(cost_function)
```


TensorBoard

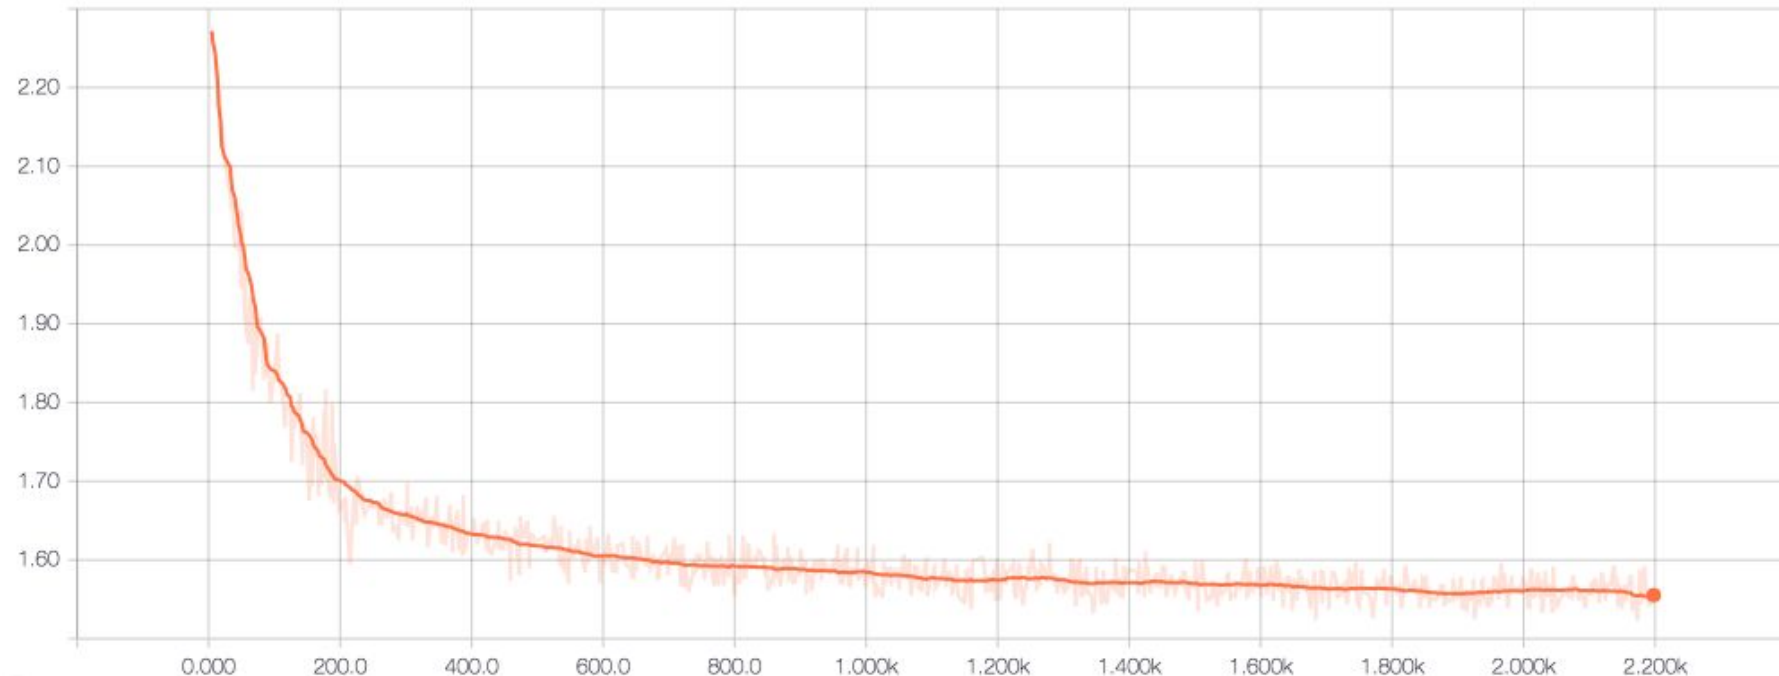


TensorBoard



Tensorboard

cost_function/cost_function



2. Multi Layer Perceptron



```
In [7]: # Set parameters
learning_rate = 0.01
training_iteration = 10
batch_size = 250
print_freq=5
```

```
In [8]: # TF graph input
x = tf.placeholder('float', [None, 784]) # mnist data image of shape 28*28=784
y = tf.placeholder('float', [None, 10]) # 0-9 digits recognition => 10 classes

keep_rate = tf.placeholder(tf.float32)
```

```
In [9]: def weight_variable(shape):
        initial = tf.constant(0.0, shape=shape)
        return tf.Variable(initial)

        def bias_variable(shape):
            initial = tf.constant(0.1, shape=shape)
            return tf.Variable(initial)
```

```
In [10]: with tf.name_scope("hidden_1") as scope:

        # Set model weights
        W_layer1 = weight_variable([784, 512])
        b_layer1 = bias_variable([512])

        # Construct a dense linear model, with act=relu and dropout
        layer_1 = tf.nn.dropout(tf.nn.relu(tf.matmul(x, W_layer1) + b_layer1), keep_rate) # Relu, dropout

        # Add summary ops to collect data
        tf.histogram_summary("W1_weights", W_layer1)
        tf.histogram_summary("B1_biases", b_layer1)
```

2. Multi Layer Perceptron



```
In [11]: with tf.name_scope("hidden_2") as scope:

    # Set model weights
    W_layer2 = weight_variable([512, 512])
    b_layer2 = bias_variable([512])

    # Construct a dense linear model, with act=relu and dropout
    layer_2 = tf.nn.dropout(tf.nn.relu(tf.matmul(layer_1, W_layer2) + b_layer2), keep_rate) # Relu, dropout

    # Add summary ops to collect data
    tf.histogram_summary("W2_weights", W_layer2)
    tf.histogram_summary("B2_biases", b_layer2)
```

```
In [12]: with tf.name_scope("output") as scope:

    # Set model weights
    W_layer3 = weight_variable([512, 10])
    b_layer3 = bias_variable([10])

    # Construct a dense linear model, with act=relu and dropout
    layer_3 = tf.add(tf.matmul(layer_2, W_layer3), b_layer3)

    # Add summary ops to collect data
    tf.histogram_summary("W3_weights", W_layer3)
    tf.histogram_summary("B3_biases", b_layer3)
```

```
In [13]: y_hat = layer_3

    # More name scopes will clean up graph representation
    with tf.name_scope("cost_function") as scope:
        # Minimize error using cross entropy
        # Cross entropy
        cost_function = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(y_hat,y))
        # Create a summary to monitor the cost function
        tf.scalar_summary("cost_function", cost_function)
```

2. Multi Layer Perceptron



```
In [14]: with tf.name_scope("train") as scope:
          # Gradient descent
          optimizer = tf.train.AdamOptimizer().minimize(cost_function)

In [15]: # Initializing the variables
          init = tf.global_variables_initializer()

          # Merge all summaries into a single operator
          merged_summary_op = tf.merge_all_summaries()

In [16]: # Launch the graph
          sess = tf.InteractiveSession()
          sess.run(init)

In [17]: # Change this to a location on your computer
          summary_writer = tf.train.SummaryWriter('./tensorboard', graph=sess.graph)

In [18]: # Training cycle
          for iteration in range(training_iteration):
              avg_cost = 0.
              total_batch = int(mnist.train.num_examples/batch_size)
              # Loop over all batches
              for i in range(total_batch):
                  batch_xs, batch_ys = mnist.train.next_batch(batch_size)

                  # dropout placeholder
                  batch_kr = 0.50

                  # Fit training using batch data
                  sess.run(optimizer, feed_dict={x: batch_xs, keep_rate: batch_kr, y: batch_ys})

                  # Compute the average loss
                  avg_cost += sess.run(cost_function, feed_dict={x: batch_xs, keep_rate: batch_kr, y: batch_ys})/(total_batch+1)

                  # Write logs for each iteration
                  summary_str = sess.run(merged_summary_op, feed_dict={x: batch_xs, keep_rate: batch_kr, y: batch_ys})
                  summary_writer.add_summary(summary_str, iteration*total_batch + i)

                  # Display logs per iteration step
                  if iteration % print_freq == 0 :
                      print("Iteration:", '%04d' % (iteration), "cost=", "{:.9f}".format(avg_cost))
```

```
Iteration: 0000 cost= 0.636337465
Iteration: 0005 cost= 0.083090350
```

2. Multi Layer Perceptron

```
In [6]: from keras.models import Sequential
        from keras.layers.core import Dense, Activation, Dropout
```

Using TensorFlow backend.

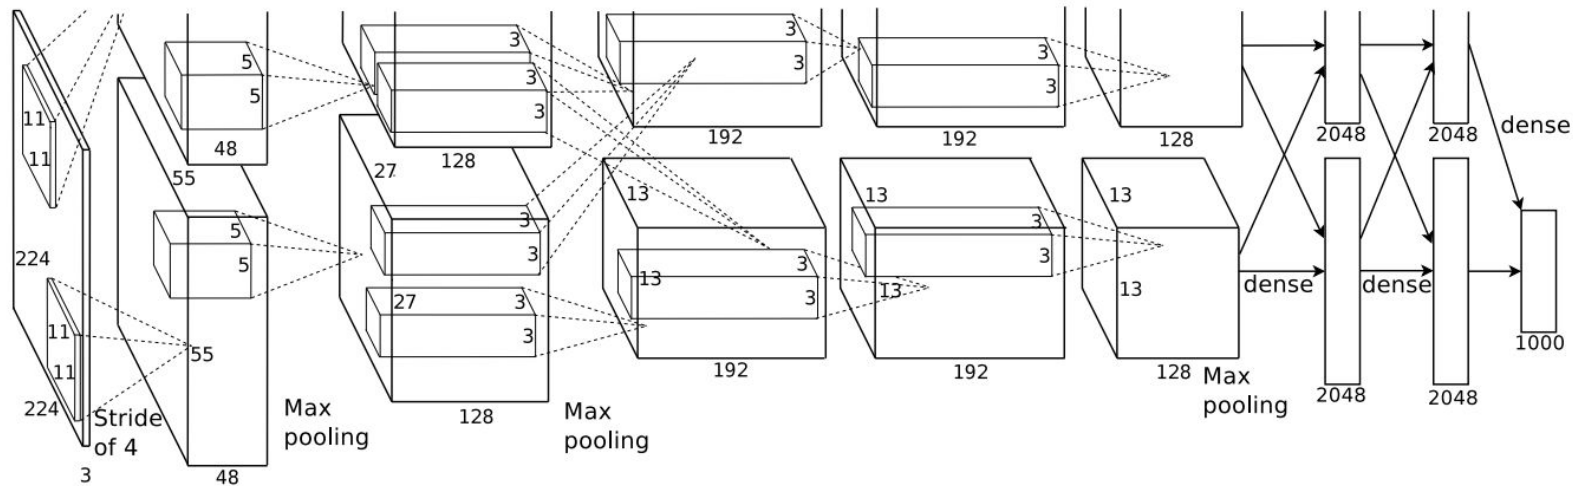
```
In [7]: model = Sequential()
        model.add(Dense(512, input_shape=(784,)))
        model.add(Activation('relu'))
        model.add(Dropout(0.25))
        model.add(Dense(512, activation='relu'))
        model.add(Activation('relu'))
        model.add(Dropout(0.25))
        model.add(Dense(10))
        model.add(Activation('softmax'))
```

```
In [8]: model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [*]: model.fit(mnist.train.images, mnist.train.labels,
                  batch_size=250, epochs=10, verbose=1,
                  validation_data=(mnist.test.images, mnist.test.labels))
```

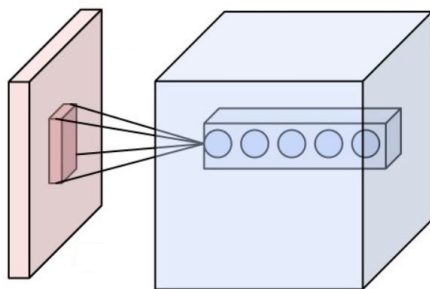


3. Convolution

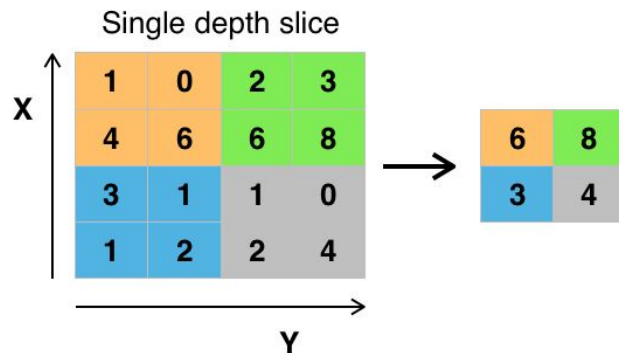


From Krizhevsky *et al.* (2012)

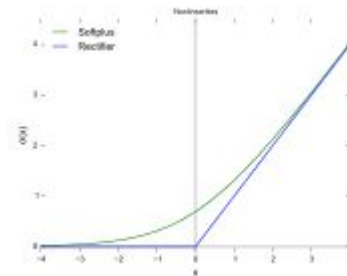
3. Convolution



convolution



Max pooling



RELU / ELU

diagrams:

By Aphex34 - Own work, CC BY-SA 4.0, <https://commons.wikimedia.org/w/index.php?curid=45659236>

By Aphex34 - Own work, CC BY-SA 4.0, <https://commons.wikimedia.org/w/index.php?curid=45673581>

CC0, <https://en.wikipedia.org/w/index.php?curid=48817276>

3. Convolution

```
In [14]: from keras.models import Sequential
         from keras.layers import Dense, Activation
         from keras.layers import Dropout, Flatten, Reshape
         from keras.layers import Conv2D, MaxPooling2D
```

```
In [15]: from keras import backend as K

         #tensorflow default channel ordering
         input_shape = (28,28,1) #channel is third
```

```
In [18]: model = Sequential()
         model.add(Reshape(input_shape, input_shape=(784,)))

         model.add(Conv2D(32, (3, 3), padding='same', activation='relu'))
         model.add(Conv2D(32, (3, 3), padding='same', activation='relu'))
         model.add(MaxPooling2D((2,2)))

         model.add(Conv2D(64, (3, 3), padding='same', activation='relu'))
         model.add(Conv2D(64, (3, 3), padding='same', activation='relu'))
         model.add(MaxPooling2D((2,2)))

         model.add(Flatten())
         model.add(Dropout(0.5))

         model.add(Dense(256, activation='relu'))
         model.add(Dropout(0.25))

         model.add(Dense(10, activation='softmax'))
```



3. Convolution

```
In [9]: from functools import reduce

for l in model.layers:
    print(l.name, l.output_shape, [reduce(lambda x, y: x*y, w.shape) for w in l.get_weights()])

reshape_1 (None, 28, 28, 1) []
convolution2d_1 (None, 28, 28, 32) [288, 32]
convolution2d_2 (None, 28, 28, 32) [9216, 32]
maxpooling2d_1 (None, 14, 14, 32) []
convolution2d_3 (None, 14, 14, 64) [18432, 64]
convolution2d_4 (None, 14, 14, 64) [36864, 64]
maxpooling2d_2 (None, 7, 7, 64) []
flatten_1 (None, 3136) []
dropout_1 (None, 3136) []
dense_1 (None, 256) [802816, 256]
dropout_2 (None, 256) []
dense_2 (None, 10) [2560, 10]
```



3. Batch Normalization

Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift

Sergey Ioffe
Google Inc., sioffe@google.com

Christian Szegedy
Google Inc., szegedy@google.com

Abstract

Training Deep Neural Networks is complicated by the fact that the distribution of each layer's inputs changes during training, as the parameters of the previous layers change. This slows down the training by requiring lower learning rates and careful parameter initialization, and makes it notoriously hard to train models with saturating nonlinearities. We refer to this phenomenon as *internal covariate shift*, and address the problem by normalizing layer inputs. Our method draws its strength from making normalization a part of the model architecture and performing the normalization for each training mini-batch. Batch Normalization allows us to use much higher learning rates and

Using mini-batches of examples, as opposed to one example at a time, is helpful in several ways. First, the gradient of the loss over a mini-batch is an estimate of the gradient over the training set, whose quality improves as the batch size increases. Second, computation over a batch can be much more efficient than m computations for individual examples, due to the parallelism afforded by the modern computing platforms.

While stochastic gradient is simple and effective, it requires careful tuning of the model hyper-parameters, specifically the learning rate used in optimization, as well as the initial values for the model parameters. The training is complicated by the fact that the inputs to each layer are affected by the parameters of all preceding layers – so that small changes to the network parameters amplify as the network becomes deeper.

[G] 2 Mar 2015

3. Batch Normalization (example for MLP)

```
In [6]: from keras.models import Model
        from keras.layers import Input, Dense, Activation
        from keras.layers import Dropout, Flatten, Reshape
        from keras.layers import Convolution2D, MaxPooling2D
        from keras.layers import BatchNormalization
```

Using TensorFlow backend.

```
In [8]: def mlp(batch_normalization=False, activation='sigmoid'):
        _in = Input(shape=(784,))

        for i in range(5):
            x = Dense(128, activation=activation, input_shape=(784,))(x if i else _in)
            if batch_normalization:
                x = BatchNormalization()(x)

        _out = Dense(10, activation='softmax')(x)
        model = Model(_in, _out)

        return model
```



3. Batch Normalization (example for MLP)

Sigmoid activation function

```
In [10]: # see http://cs231n.github.io/neural-networks-3/
```

```
model = mlp(False, 'sigmoid')
print_layers(model)

bl_noBN = BatchLogger()

from keras.optimizers import Adam
model.compile(loss='categorical_crossentropy', optimizer=Adam(lr=0.001), metrics=["accuracy"])

model.fit(mnist.train.images, mnist.train.labels,
          batch_size=128, nb_epoch=1, verbose=1, callbacks=[bl_noBN],
          validation_data=(mnist.test.images, mnist.test.labels))
```

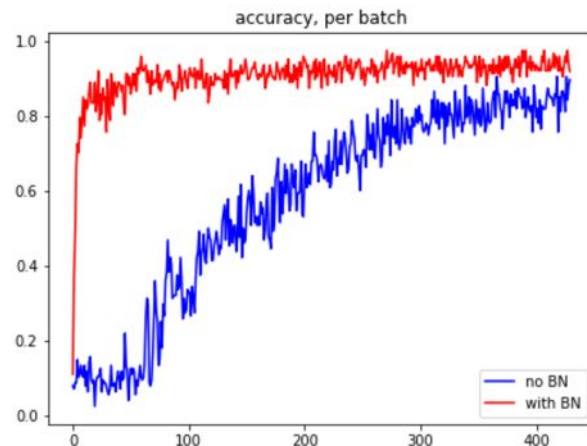
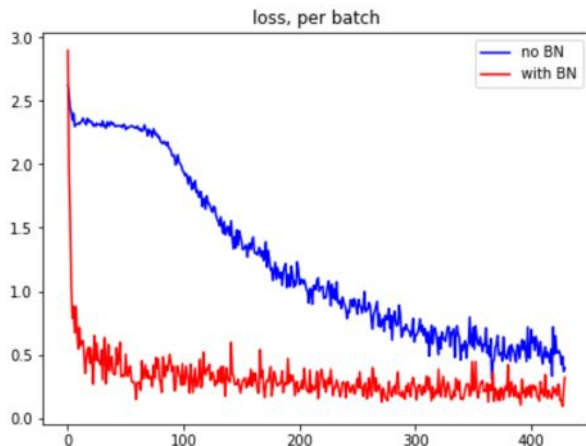
```
input_1 (None, 784) []
dense_1 (None, 128) [100352, 128]
dense_2 (None, 128) [16384, 128]
dense_3 (None, 128) [16384, 128]
dense_4 (None, 128) [16384, 128]
dense_5 (None, 128) [16384, 128]
dense_6 (None, 10) [1280, 10]
```



3. Batch Normalization (example for MLP)

```
In [12]: plt.figure(figsize=(15,5))
plt.subplot(1, 2, 1)
plt.title('loss, per batch')
plt.plot(bl_noBN.log_values['loss'], 'b-', label='no BN');
plt.plot(bl_BN.log_values['loss'], 'r-', label='with BN');
plt.legend(loc='upper right')
plt.subplot(1, 2, 2)
plt.title('accuracy, per batch')
plt.plot(bl_noBN.log_values['acc'], 'b-', label='no BN');
plt.plot(bl_BN.log_values['acc'], 'r-', label='with BN');
plt.legend(loc='lower right')
plt.show()
```

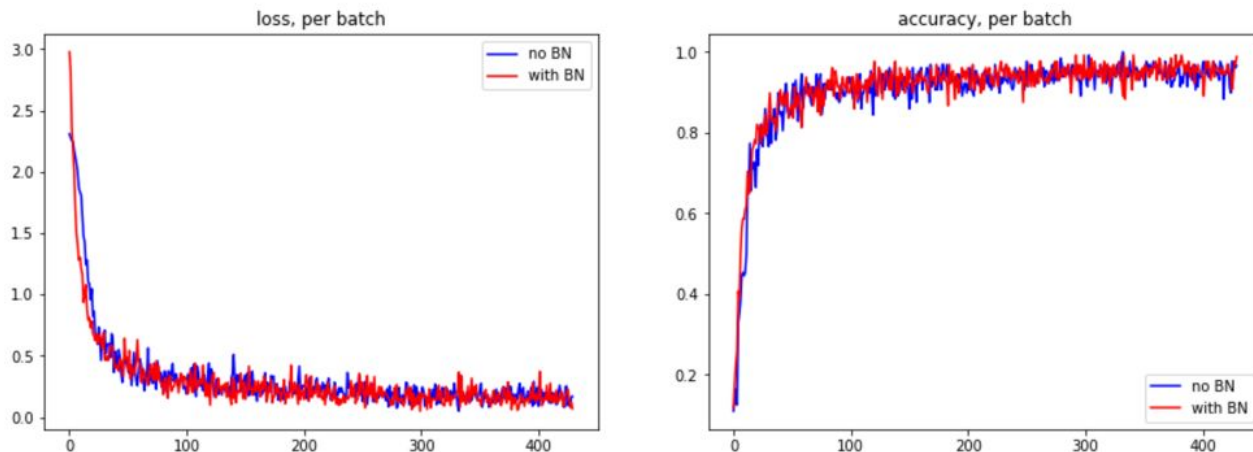
Activation function:
SIGMOID



3. Batch Normalization (example for MLP)

```
In [15]: plt.figure(figsize=(15,5))
plt.subplot(1, 2, 1)
plt.title('loss, per batch')
plt.plot(bl_noBN.log_values['loss'], 'b-', label='no BN');
plt.plot(bl_BN.log_values['loss'], 'r-', label='with BN');
plt.legend(loc='upper right')
plt.subplot(1, 2, 2)
plt.title('accuracy, per batch')
plt.plot(bl_noBN.log_values['acc'], 'b-', label='no BN');
plt.plot(bl_BN.log_values['acc'], 'r-', label='with BN');
plt.legend(loc='lower right')
plt.show()
```

Activation function:
RELU

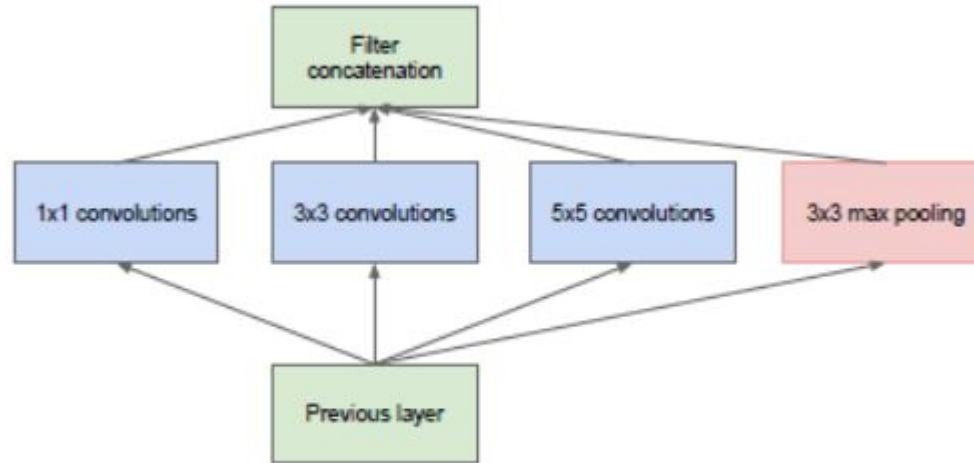


4. Regularization: Prevent overfitting in ANNs

- Batch Normalization
- RELU/ELU
- RESIDUAL / SKIP Networks
- DROP LAYER
- REDUCE PRECISION (HUFFMAN ENCODING)

In general ANN are parameters rich, constraining the parameter space usually produces better results and speed up the learning

5. Inception architectures



(a) Inception module, naïve version

Cannot be stacked!

5. Inception architectures

```
In [15]: def BNConv(filters, nb_row, nb_col, subsample=(1, 1), padding="same"):
    def f(input):
        conv = Conv2D(activation="relu", kernel_size=(nb_row, nb_col), filters=filters, strides=
            padding=padding, kernel_initializer="he_normal")(input)
        return BatchNormalization()(conv)
    return f
```

```
In [16]: def inception_naive_module(m=1):
    def f(input):

        # Tower A
        conv_a = BNConv(32*m, 1, 1)(input)

        # Tower B
        conv_b = BNConv(32*m, 3, 3)(input)

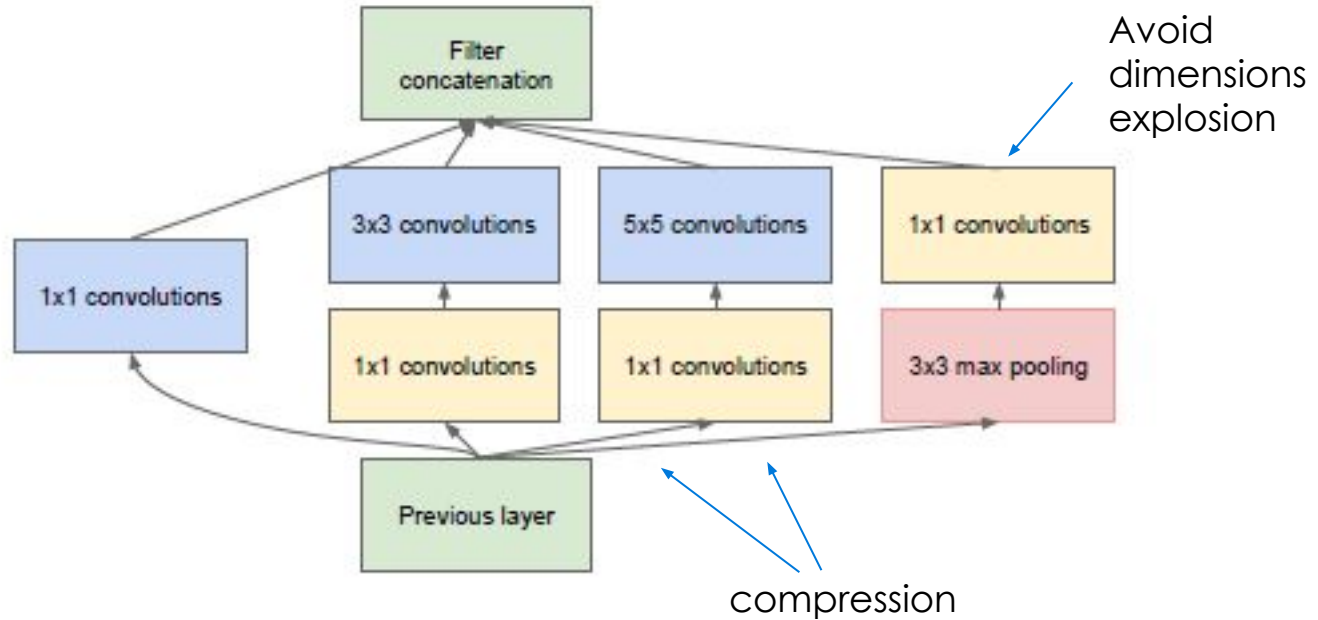
        # Tower C
        conv_c = BNConv(16*m, 5, 5)(input)

        # Tower D
        pool_d = MaxPooling2D(pool_size=(3, 3), strides=(1, 1), padding="same")(input)

        return merge([conv_a, conv_b, conv_c, pool_d], mode='concat', concat_axis=3)

    return f
```

5. Inception architectures



5. Inception architectures

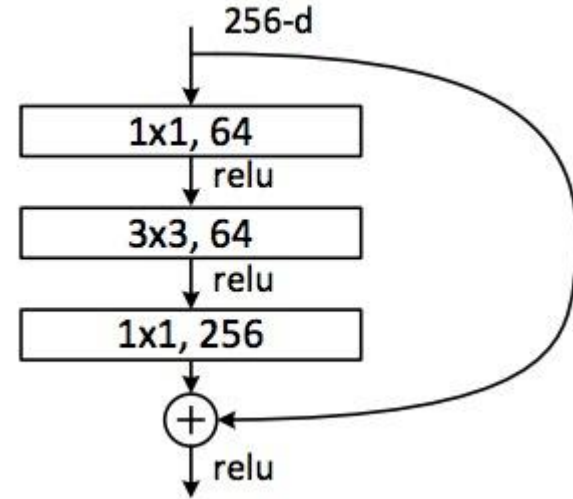
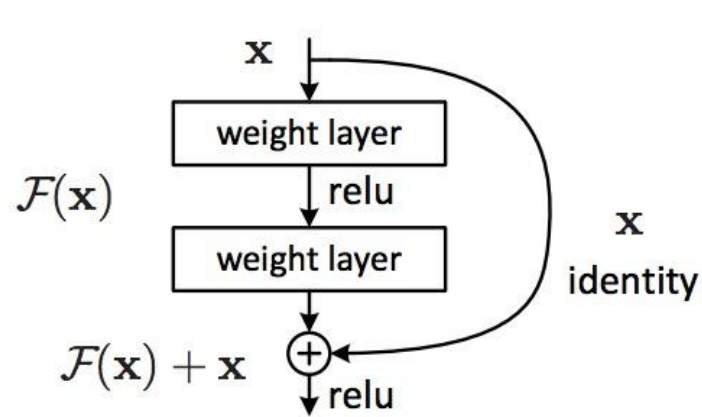
```
In [23]: def inception_dimred_module(m=1):  
    def f(input):  
  
        # Tower A  
        conv_a = BNConv(32*m, 1, 1)(input)  
  
        # Tower B  
        conv_b = BNConv(16*m, 1, 1)(input)  
        conv_b = BNConv(32*m, 3, 3)(conv_b)  
  
        # Tower C  
        conv_c = BNConv(4*m, 1, 1)(input)  
        conv_c = BNConv(16*m, 5, 5)(conv_c)  
  
        # Tower D  
        # max pooling followed by compression  
        pool_d = MaxPooling2D(pool_size=(3, 3), strides=(1, 1), padding="same")(input)  
        conv_d = BNConv(16*m, 1, 1)(pool_d)  
  
        return merge([conv_a, conv_b, conv_c, conv_d], mode='concat', concat_axis=3)  
  
    return f
```

5. Inception architectures (top architecture)

```
In [49]: def inception(inception_module):  
  
    #input in the right shape, tensorflow ordered  
    _in = Input(shape=(784,))  
    reshape_1 = Reshape((28,28,1))(_in)  
  
    # go to 32 channels  
    conv_0 = BNConv(32, 3, 3)(reshape_1)  
    conv_0 = BNConv(32, 3, 3)(conv_0)  
    pool_0 = MaxPooling2D((2, 2))(conv_0)  
  
    # apply inception network (input: 14x14x32, output channels:96)  
    module_1 = inception_module()(pool_0)  
  
    # pool to 7x7x96  
    pool_1 = MaxPooling2D((2, 2))(module_1)  
  
    # apply inception network (input: 7x7x96, output channels:192)  
    module_2 = inception_module(m=2)(pool_1)  
  
    # pool to: 1x1x96 and flatten  
    x = AveragePooling2D((7, 7))(module_2)  
    x = Flatten()(x)  
    x = Dropout(0.4)(x)  
  
    # dense layer and normalization  
    fc = Dense(128, activation='relu')(x)  
    fc = BatchNormalization()(fc)  
  
    _out = Dense(10, activation='softmax')(fc)  
    model = Model(_in, _out)  
  
    return model
```



6. Residual Networks



<https://culurciello.github.io/tech/2016/06/04/nets.html>

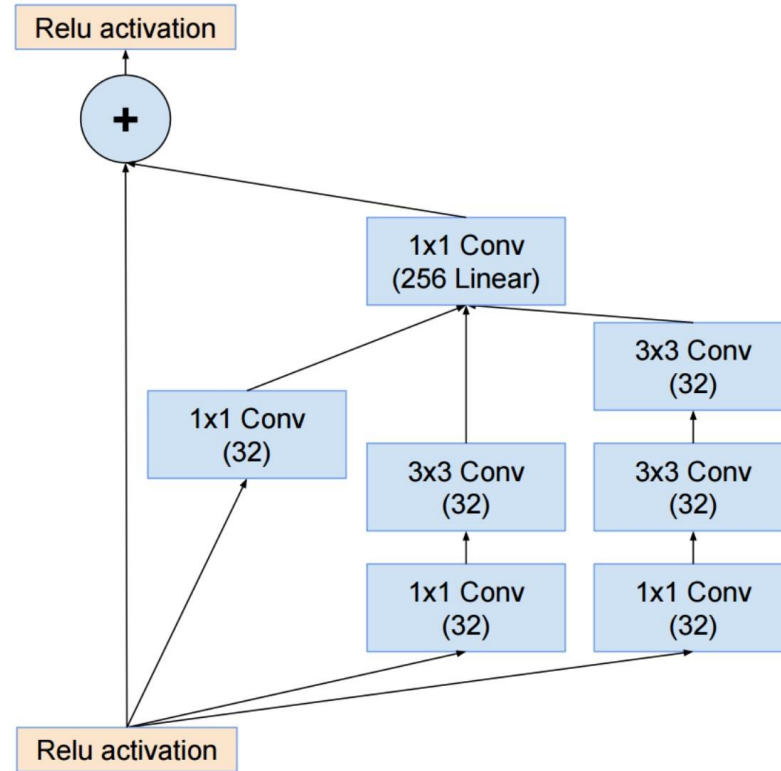
6. Residual Networks

```
In [8]: def residual_block(skip=True):  
        def f(input):  
            conv = Convolution2D(4,3,3,border_mode='same', activation='relu')(input)  
            res = merge([conv,input], mode='sum')  
            return Activation('relu')(res) if skip else conv  
        return f
```

```
In [9]: def resnet(skiplayers=3):  
        #select inception module  
        _in = Input(shape=(784,))  
        reshape = Reshape((28,28,1))(_in)  
        res = Convolution2D(4,3,3,border_mode='same')(reshape)  
  
        for i in range(skiplayers):  
            res = residual_block()(res)  
  
        flat = Flatten()(res)  
        flat = Dropout(0.4)(flat)  
  
        _out = Dense(10, activation='softmax')(flat)  
        model = Model(_in, _out)  
  
        return model
```

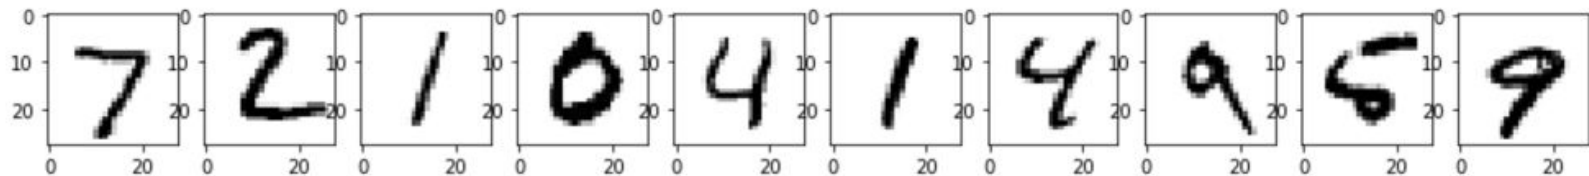
```
In [10]: model = resnet(10)
```


6. Residual + Inception Networks



7. LSTM on Images

```
In [5]: plt.figure(figsize=(15,5))
for i in list(range(10)):
    plt.subplot(1, 10, i+1)
    pixels = mnist.test.images[i]
    pixels = pixels.reshape((28, 28))
    plt.imshow(pixels, cmap='gray_r')
plt.show()
```



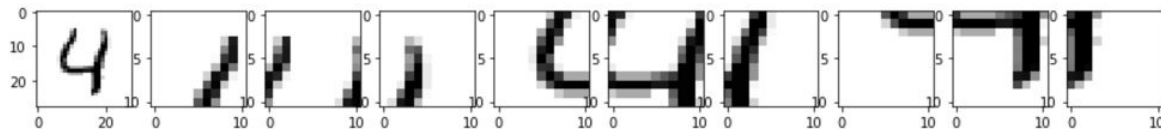
7. LSTM on Images

```
In [7]: import math
def glimpses(pixels, n=1):
    g = []
    k = int(math.sqrt(n-1))//2
    r = list(range(-k, k+1))
    if type(n)==list:
        r = n
    for i in r:
        for j in r:
            g.append(glimpse(pixels,14+7*j,14+7*i,5))
    return np.array(g)
```

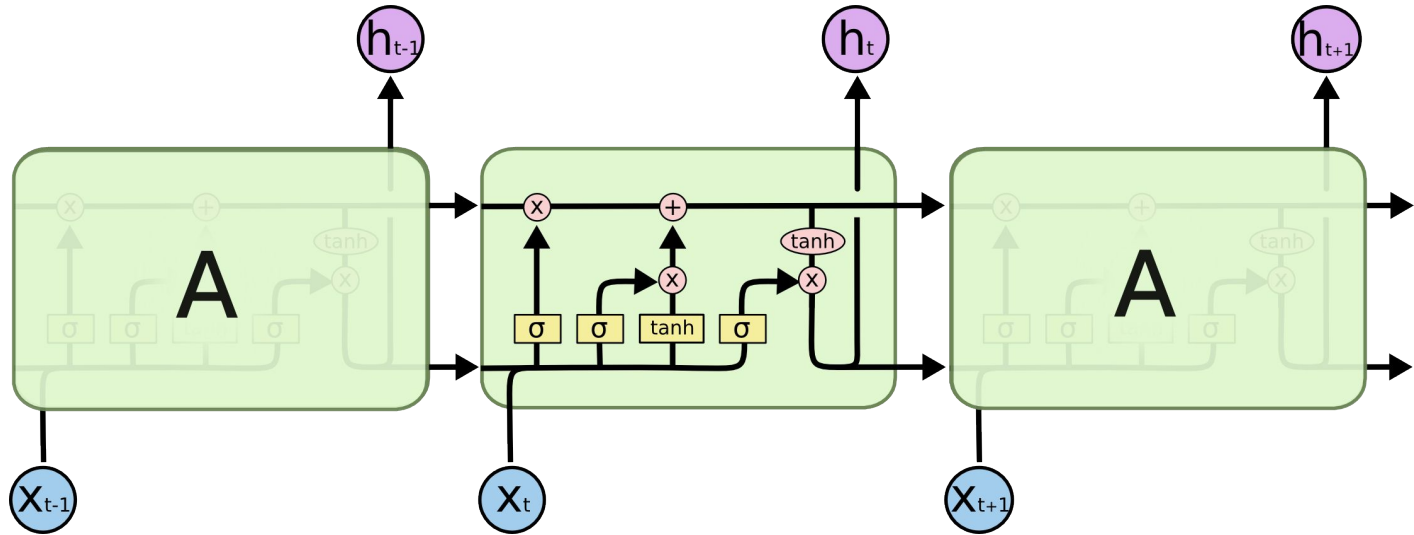
```
In [8]: pixels = mnist.test.images[4]
plt.figure(figsize=(15,5))

# plot the full field
plt.subplot(1, 10, 1)
plt.imshow(glimpse(pixels,14,14,14), cmap='gray_r')

# plot 9 glimpses
i = 2
for g in glimpses(pixels,9):
    plt.subplot(1, 10, i)
    plt.imshow(g, cmap='gray_r')
    i += 1
plt.show()
```

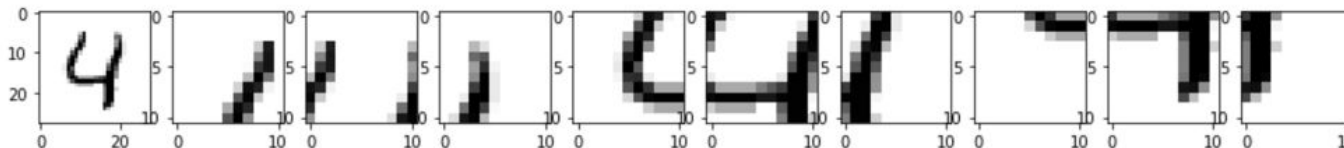


7. LSTM on Images



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

7. LSTM on Images



```
In [144]: from keras.models import Sequential
          from keras.layers import Dense
          from keras.layers import LSTM

          model = Sequential()
          model.add(LSTM(32, input_length=9, input_dim=121))
          model.add(Dense(10, activation='softmax'))

          from keras.optimizers import Adam
          model.compile(loss='categorical_crossentropy', optimizer=Adam(lr=0.01), metrics=["accuracy"])
```

```
In [145]: from keras.utils.layer_utils import print_summary
          print_summary(model.layers)
```



7. LSTM on Images

```
In [146]: def sequence(x,n):  
          return glimpses(x,n).reshape((n*121))
```

```
In [147]: # prepare the train/test input as a tensor of shape 55000, 9, 121  
train_sequences = np.apply_along_axis(sequence, 1, mnist.train.images,9).reshape(-1,9,121)  
test_sequences  = np.apply_along_axis(sequence, 1, mnist.test.images,9).reshape(-1,9,121)
```

```
In [148]: model.fit(train_sequences, mnist.train.labels,  
                    batch_size=250, nb_epoch=5, verbose=1,  
                    validation_data=(test_sequences, mnist.test.labels))
```



7. LSTM on ConvNets (bonus slide)

```
In [215]: # Convolutional + Multilayer LSTM
from keras.models import Model

from keras.layers import TimeDistributed
from keras.layers import Flatten, Reshape
from keras.layers import Convolution2D, MaxPooling2D, BatchNormalization

def conv_net():
    def f(_in):
        # go to 32 channels
        layer = Convolution2D(16, 3, 3, border_mode="same", activation="relu")(_in)
        layer = Convolution2D(16, 3, 3, border_mode="same", activation="relu")(_in)
        layer = MaxPooling2D((2, 2))(layer)
        layer = Flatten()(layer)
        _out = Dense(64, activation='relu')(layer)
        return _out
    return f

# create the conv_model
x = Input(shape=(11, 11, 1))
conv_model = Model(x, conv_net()(x))

# build the top model
model = Sequential()

# prep for convolution, keep the timestep as first dimension (after the implicit batch dim)
model.add(Reshape((9,11,11,1), input_shape=(9,121)))

# time distributed on the convolutional part
model.add(TimeDistributed(conv_model))

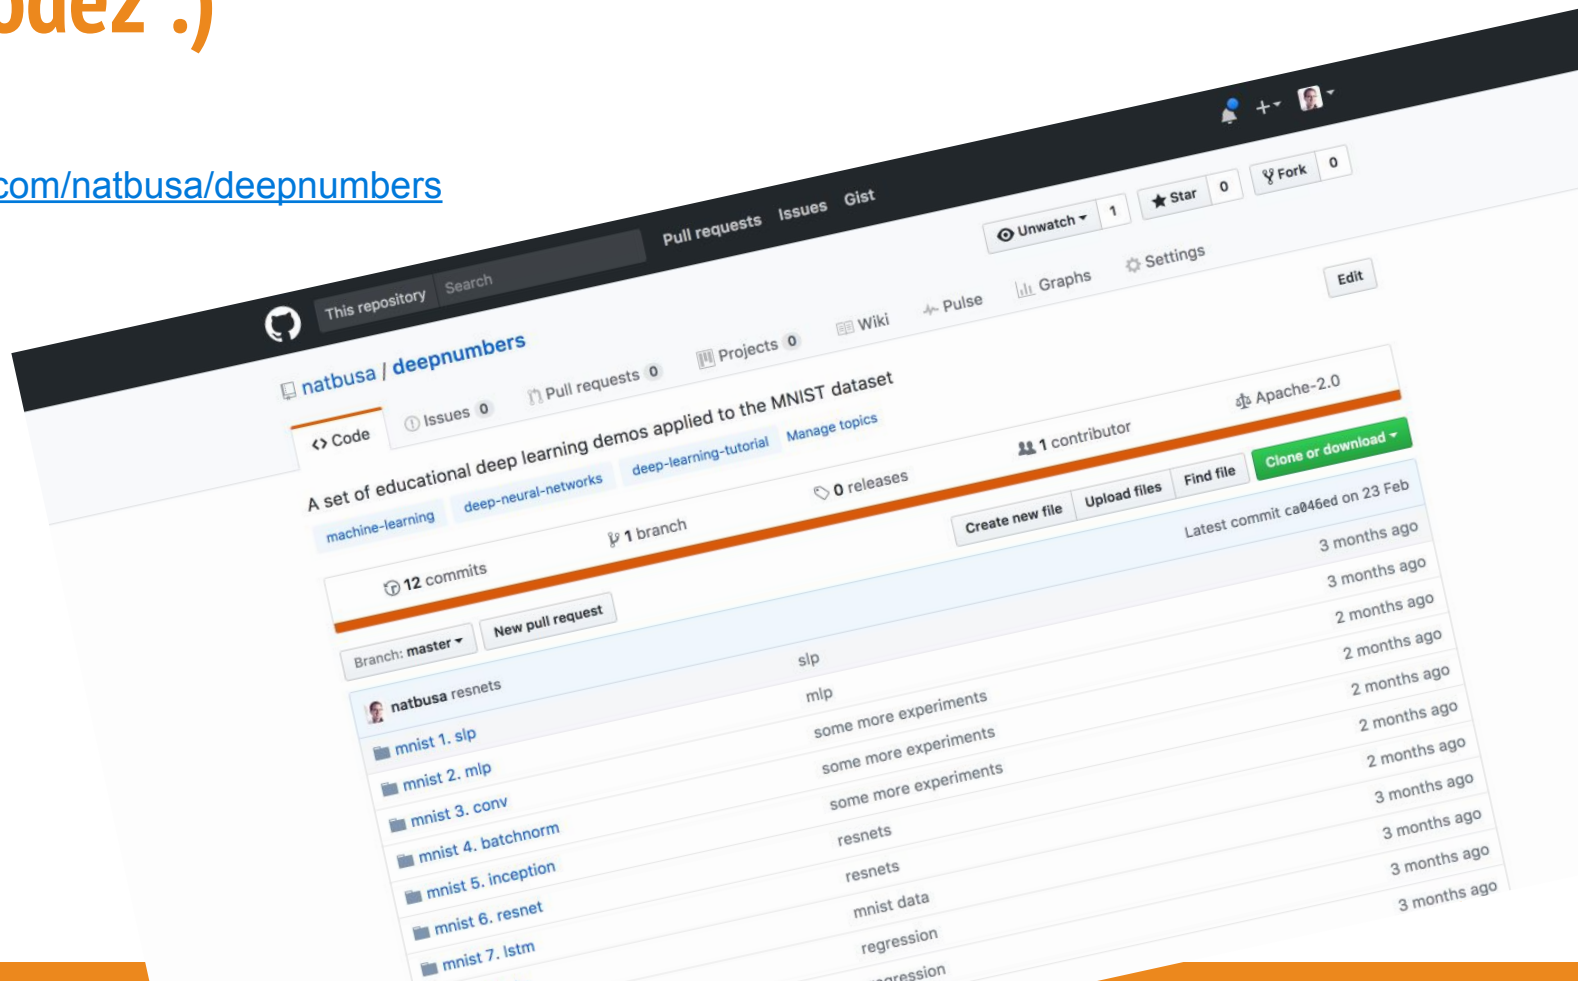
# temporal model (64 is the output dim of the conv_model)
model.add(LSTM(32, return_sequences=True, input_length=9, input_dim=64))
model.add(LSTM(32, return_sequences=True))
model.add(LSTM(32))

# last layer
model.add(Dense(10, activation='softmax'))

from keras.optimizers import Adam
model.compile(loss='categorical_crossentropy', optimizer=Adam(lr=0.01), metrics=["accuracy"])
```

All the codez :)

<https://github.com/natbusa/deepnumbers>



Meta- References

... just a few articles, but extremely dense in content.
A must read!

<https://keras.io/>

<http://karpathy.github.io/neuralnets/>

<https://culurciello.github.io/tech/2016/06/04/nets.html>

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

<https://gab41.lab41.org/batch-normalization-what-the-hey-d480039a9e3b>